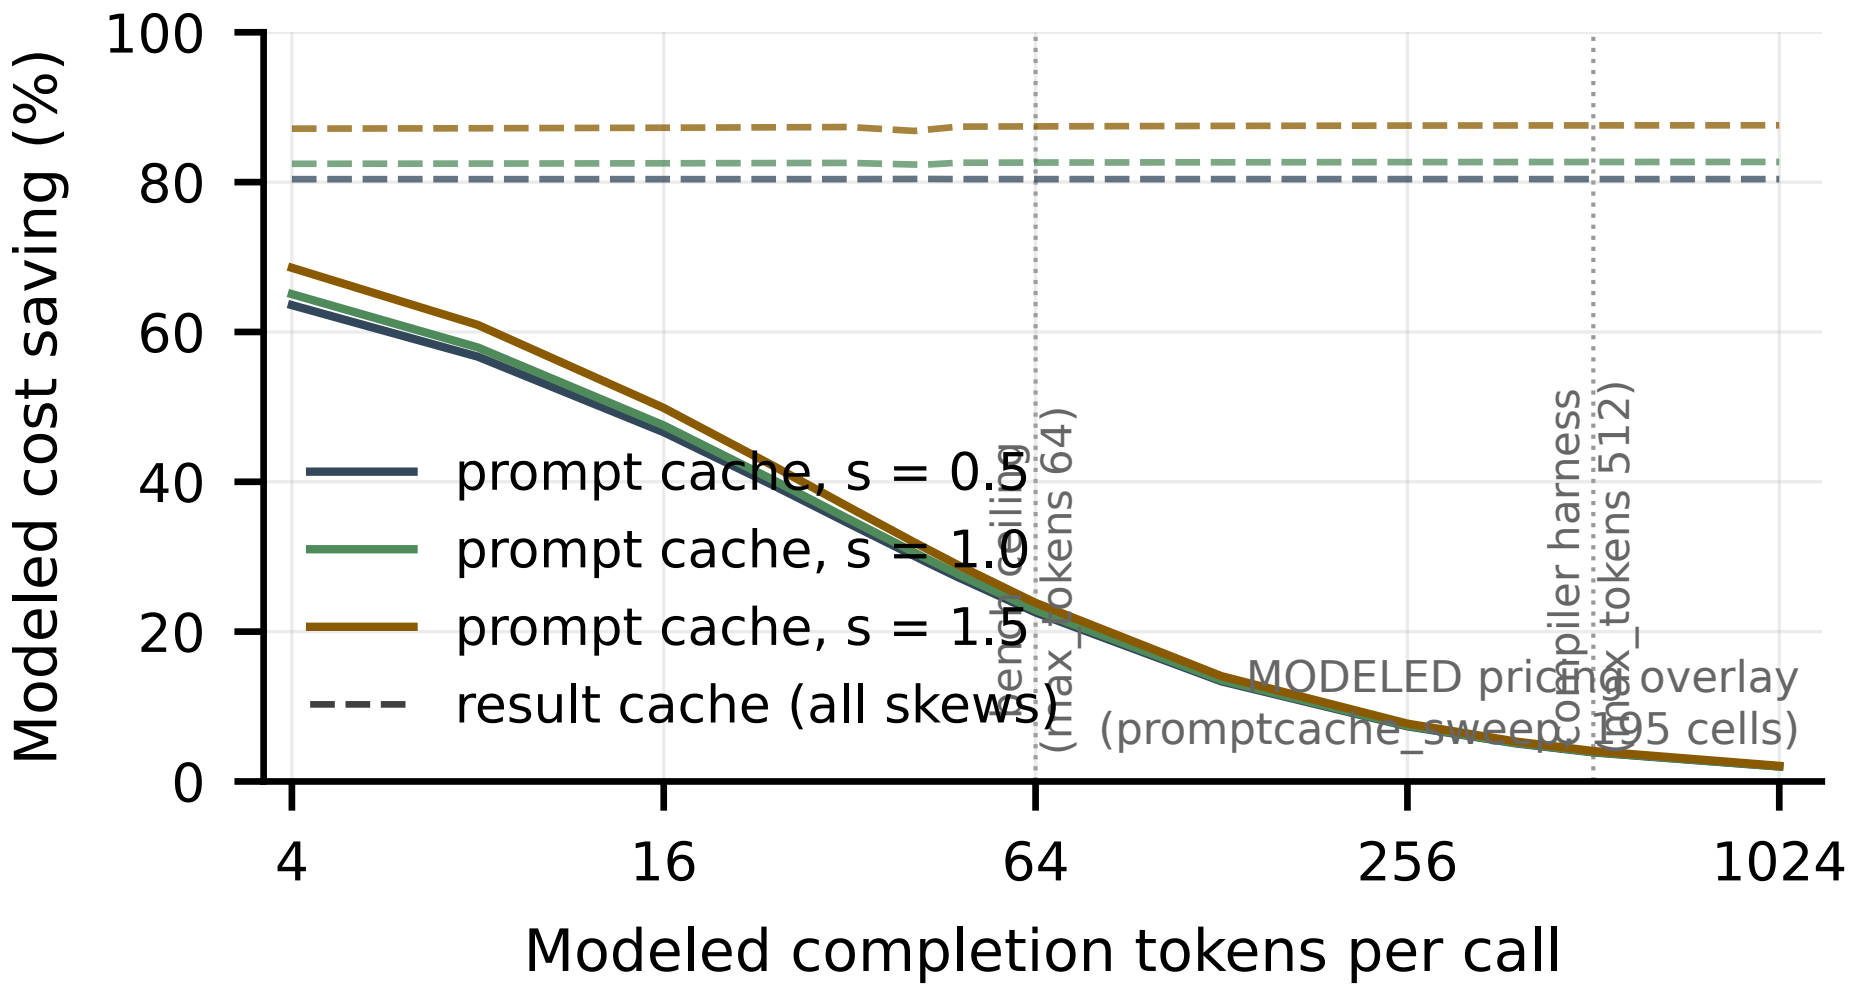
Modeled cost saving (%)
100
80
60
40
20
0
prompt cache, s = 0.5
prompt cache, s = 1.0
prompt cache, s = 1.5
result cache (all skews)
typical ceiling
(max_tokens 64)
compiler harness
(max_tokens 512)
MODELED pricing overlay
(promptcache_sweep, 195 cells)
4
16
64
256
1024
Modeled completion tokens per call